



AR-021 Confidence: 82%

AI Agent Economics The Real ROI Nobody Talks About

Why Enterprise AI Agent Deployments Cost 3-7x Initial Estimates

February 2026

v1.0

Florian Ziesche · Ainary Ventures

CONTENTS

FOUNDATION

| | | |
|---|-------------------------|---|
| 1 | How to Read This Report | 3 |
| 2 | Executive Summary | 4 |
| 3 | Methodology | 5 |

ANALYSIS

| | | |
|---|---|----|
| 4 | The Vendor Claim: "\$0.50 Per Interaction" | 6 |
| 5 | The Hidden Costs Enterprise Teams Discover | 8 |
| 6 | Trust Infrastructure: The Multiplier Nobody Budgets For | 11 |
| 7 | The Economics of Agent Failure | 14 |
| 8 | Real-World Economics: What We Learned Building 18 Reports | 16 |
| 9 | When Agent ROI Actually Works | 19 |

ACTION

| | | |
|----|-------------------|----|
| 10 | Recommendations | 22 |
| 11 | Transparency Note | 24 |
| 12 | Claim Register | 25 |

1. How to Read This Report

This report uses a structured confidence rating system to communicate what is known versus what is inferred. Every quantitative claim carries its source and confidence level.

| RATING | MEANING | EXAMPLE |
|--------|--|---|
| High | 3+ independent sources, verifiable data | Maintenance costs 15–30% of development costs annually (industry surveys) |
| Medium | 1–2 sources, plausible but not independently confirmed | Security adds 20–40% to platform costs (vendor estimates) |
| Low | Single secondary source, methodology unclear | 89% of agents never reach production (single vendor claim) |

This report was produced using a **multi-agent research pipeline** with cost tracking and structured source validation. Full methodology details are provided in the Transparency Note (Section 11).

2. Executive Summary

Most AI agent ROI calculations are fantasy. Actual deployment costs are 3-7x initial estimates when you include trust infrastructure, monitoring, failure remediation, and the organizational overhead nobody budgets for.

- **Vendor pitch: "\$0.25–\$0.50 per interaction"** ignores setup costs (\$100K–\$2M), ongoing monitoring (15–30% of dev costs annually), security compliance (20–40% of platform costs), and failure remediation^{[1][2][3]}
- **Trust infrastructure alone adds 1.5x to agent deployment costs** through observability, human oversight, audit trails, and rollback mechanisms — costs that don't exist in chatbot deployments^[4]
- **Failure correction costs exceed prevention costs** when AI workflow error rate hits 10% — requiring Tier 3 intervention and reversibility infrastructure that most teams discover mid-deployment^[5]
- **89% of enterprise AI agents never reach production**, making the real cost question "What did we spend to learn it won't work?" not "How much did we save?"^[6]
- **Token optimization and model selection can reduce costs 40–50%** — proven in Ainary's own production: 18 research reports at \$2.75 average API cost through Sonnet vs Opus switching and aggressive caching^[Internal]

Keywords: *AI agent economics, total cost of ownership, trust infrastructure, failure remediation, ROI calculation, agent deployment costs, hidden costs*

3. Methodology

This report synthesizes enterprise deployment data from vendor case studies, industry surveys (Deloitte, Gartner, PwC), and direct cost tracking from Ainary's own agent production environment (18 published research reports). We analyzed 14 vendor pricing models, 6 enterprise case studies with disclosed costs, and 3 cost estimation frameworks published in 2025-2026.

Limitations: Few enterprises publicly disclose AI agent deployment costs. Most cost data comes from vendor-published case studies (selection bias toward successful deployments) or practitioner surveys with limited sample sizes. Our "3-7x initial estimate" range is directionally supported but not rigorously quantified across a large sample.

Full methodology details, including confidence calibration and known weaknesses, are provided in the Transparency Note (Section 11).

4. The Vendor Claim: "\$0.50 Per Interaction"

85%

(Confidence: High)

The industry-standard cost comparison — AI at \$0.25–\$0.50 per interaction vs. human agents at \$3–\$8 — is accurate but incomplete. It measures marginal cost per transaction while ignoring every cost that makes those transactions possible.

The Math Is Correct, The Framing Is Deceptive

Multiple independent sources confirm the per-interaction cost advantage^{[1][7]}:

- **AI agents:** \$0.25–\$0.50 per customer service interaction
- **Human agents:** \$3.00–\$6.00 (mid-tier markets), \$4.00–\$8.00 (Tier 1 markets like NA, UK)
- **Cost reduction:** 85–90% on a per-interaction basis

This comparison is not wrong — it is just radically incomplete. It is the equivalent of comparing the cost of sending an email (\$0.00) to the cost of mailing a letter (\$0.68) while ignoring the cost of the email server, the internet connection, the IT team, and the spam filter.

Exhibit 1: Vendor Cost Comparison vs. Total Cost of Ownership

| COST COMPONENT | INCLUDED IN VENDOR PITCH | ACTUAL ENTERPRISE COST |
|----------------------------|--------------------------|--|
| Per-interaction API cost | Yes (\$0.25–\$0.50) | Yes (\$0.25–\$0.50) |
| Setup / integration | No | \$100K–\$2M |
| Ongoing monitoring | No | 15–30% of dev costs annually |
| Security & compliance | No | 20–40% of platform costs |
| Human oversight (HITL) | No | 1 FTE per 3–5 agents (\$80K–\$150K/year) |
| Failure remediation | No | Variable, often exceeds prevention costs |
| Model retraining / updates | No | \$10K–\$50K per cycle |
| Change management | No | Often the largest cost (not quantified) |

Source: Analysis of vendor pricing (Teneo.ai, Riseup Labs, Hypersense Software) [1][2][3] vs. enterprise TCO studies [4][6][8]

When enterprises budget based on per-interaction cost, they allocate for the API bill. When reality hits, they discover they also need to pay for everything that makes those API calls trustworthy, auditable, and reversible.

WHAT WOULD INVALIDATE THIS?

If a vendor emerged with a fully managed AI agent service that included trust infrastructure, monitoring, and compliance in a transparent all-in price, the cost gap would narrow. Most vendors charge for the model but not the infrastructure to make it safe.

SO WHAT?

When evaluating agent ROI, demand Total Cost of Ownership breakdowns. Ask vendors: "What costs are not in this quote that I will discover in months 3–12?" The honest vendors will tell you. The dishonest ones will discover it with you.

5. The Hidden Costs Enterprise Teams Discover

78%

(Confidence: High)

The cost components nobody budgets for in year 1 become the largest line items in year 2.

1. Setup and Integration: The One-Time Cost That Never Ends

Initial deployment costs vary wildly based on complexity^{[2][3][8]}:

- **Basic agent (single function, no memory):** \$5K–\$15K (prompt engineering only)
- **Mid-complexity (RAG, tools, basic orchestration):** \$100K–\$200K
- **Enterprise multi-agent system:** \$500K–\$2M+

The "one-time" framing is deceptive. Integration is never finished because:

- APIs change (model upgrades, deprecations, pricing shifts)
- Business requirements evolve faster than agent capabilities
- Legacy system integrations break when either side updates
- Security requirements tighten post-deployment (see: CVE-2025-32711 EchoLeak forcing emergency upgrades^[9])

2. Monitoring and Observability: 15–30% of Dev Costs Annually

Enterprise-grade agent monitoring requires^{[2][10]}:

- **Observability platforms:** LangSmith, Arize, Langfuse (\$5K–\$50K/year depending on scale)
- **Custom instrumentation:** Logging, tracing, performance metrics (engineering time)
- **Alert infrastructure:** Anomaly detection, threshold monitoring, escalation workflows

- **Dashboard maintenance:** Continuous updates as agent behavior evolves

Industry estimates: **15–30% of initial development costs** annually just to maintain visibility into what agents are doing^[2]. This is not a luxury — it is the minimum requirement to know when an agent starts failing before customers report it.

3. Security and Compliance: 20–40% Platform Cost Adder

Compliance and security are non-negotiable for enterprise deployments^[3]:

- **Security audits:** Third-party penetration testing, vulnerability assessments
- **Compliance certification:** SOC 2, ISO 42001, EU AI Act readiness
- **Ongoing monitoring:** Intrusion detection, access logging, anomaly detection
- **Regulatory overhead:** Legal review, risk assessments, documentation

Typical adder: **20–40% of platform costs** for organizations in regulated industries^[3]. For high-risk deployments (healthcare, financial services), this can exceed 50%.

4. Human Oversight: The "Autonomous" Agent That Needs Babysitting

Most production agents require continuous human oversight^[4]:

- **Tier 1 oversight:** Monitoring dashboards, reviewing flagged interactions (\$50K–\$80K/year per FTE)
- **Tier 2 intervention:** Correcting agent errors, updating prompts, retraining workflows (\$80K–\$120K/year per FTE)
- **Tier 3 escalation:** Handling edge cases the agent cannot resolve, crisis management (\$120K–\$200K/year per senior engineer)

Rule of thumb from enterprise deployments: **1 FTE per 3–5 production agents** for ongoing oversight. This is not included in per-interaction cost calculations.

5. Change Management: The Silent Budget Killer

Technical readiness is only 30% of successful deployment^[11]. The other 70% is organizational change management:

- Training teams to work with agents (not just use them)
- Redesigning workflows around agent capabilities
- Managing employee anxiety about displacement
- Updating policies, procedures, and documentation
- Building trust in agent outputs through repeated validation

McKinsey and BCG research on digital transformation suggests change management costs **exceed technology costs by 2-3x** in complex organizations^[11]. Most AI agent budgets allocate zero dollars for this.

WHAT WOULD INVALIDATE THIS?

If a class of "truly autonomous" agents emerged that required zero human oversight, operated within perfect security boundaries, and integrated seamlessly with legacy systems, these hidden costs would disappear. No such agent exists today.

SO WHAT?

Budget for 1.5x your initial cost estimate in year 1, then 30–50% of dev costs annually for maintenance. If that destroys your ROI, the deployment is not viable. Better to learn this before spending \$500K than after.

6. Trust Infrastructure: The Multiplier Nobody Budgets For

75%

(*Confidence: Medium — based on directional evidence, not rigorous cost accounting*)

Trust infrastructure is the difference between a chatbot and an enterprise agent. It is also the cost multiplier that turns "\$100K deployment" into "\$350K reality."

What Is Trust Infrastructure?

Trust infrastructure is everything required to make an AI agent's outputs trustworthy in a high-stakes environment^[4]:

- **Provenance tracking:** Where did this data come from? (see AR-014 Agent Memory, AR-015 Knowledge Compounding)
- **Confidence scoring:** How certain is the agent? (see AR-009 Calibration Gap)
- **Audit trails:** What did the agent do and why?
- **Rollback mechanisms:** Can we undo this action if it was wrong?
- **Human-in-the-loop (HITL) workflows:** When should a human decide? (see AR-011 HITL Illusion)
- **Escalation paths:** What happens when the agent fails?
- **Blast radius containment:** How do we prevent one bad agent action from cascading?

Chatbots do not need this infrastructure — if a chatbot gives a wrong answer, the user ignores it or asks again. Agents that take actions (send emails, place orders, modify databases, execute code) need every component listed above.

Exhibit 2: Chatbot vs. Agent — Cost of Trust Infrastructure

| TRUST COMPONENT | CHATBOT (STATIC Q&A) | AGENT (ACTIONS + MEMORY) |
|--------------------------------|----------------------|--|
| Provenance tracking | Not required | Required (\$20K–\$60K setup) |
| Audit trails (immutable) | Nice to have | Compliance requirement (\$15K–\$40K) |
| Rollback / undo infrastructure | Not applicable | Critical (\$30K–\$100K) |
| HITL escalation workflows | Optional | Required (\$10K–\$30K) |
| Confidence calibration | Optional | Required for high-stakes decisions (\$20K–\$50K) |
| Multi-agent coordination | Not applicable | Required for complex workflows (\$40K–\$120K) |
| Total Trust Tax | \$0–\$20K | \$135K–\$400K |

Source: Author estimates based on AR-002 Trust Tax, enterprise deployment case studies [4][8], and vendor pricing for observability/governance tools

The trust infrastructure cost is why **agent deployments cost 1.5x chatbot deployments** even when using the same underlying LLM. You are not just paying for the model — you are paying for the safety net that makes the model safe to deploy.

The Klarna Warning: When You Skip Trust Infrastructure

Klarna deployed AI agents aggressively in 2024-2025, replacing 853 FTEs and saving an estimated \$60M^[12]. By Q4 2025, CEO Sebastian Siemiatkowski admitted they had "overpivoted" to AI, causing quality problems that required course correction^[13].

The lesson: **Klarna optimized for cost reduction without proportional investment in trust infrastructure.** When agent quality degraded, they

discovered the cost of fixing errors exceeded the cost of preventing them.

The economic principle: **Correction costs exceed prevention costs when error rates hit 10%**. If an AI agent has a 90% success rate and the 10% failures require expensive Tier 3 intervention, the ROI can turn negative^[5]. Trust infrastructure is the investment that keeps error rates below the break-even threshold.

CLAIM

Trust infrastructure adds 1.5x to agent deployment costs but prevents the 3x correction costs that destroy ROI when agents fail at scale.

WHAT WOULD INVALIDATE THIS?

If models achieved 99.9% reliability in production without trust infrastructure, the entire "trust tax" thesis would collapse. Current evidence: most production agents operate at 85–95% success rates, making trust infrastructure non-optional.

SO WHAT?

Do not deploy agents without budgeting for trust infrastructure. The cost is real, but the alternative — fixing production failures at scale — is always more expensive. See AR-002 Trust Tax for the full economic framework.

7. The Economics of Agent Failure

70%

(Confidence: Medium — directional evidence strong, quantification estimates vary)

The cost of an AI agent failure is not the cost of the API call. It is the cost of detecting the failure, diagnosing it, correcting it, and preventing it from happening again.

The Correction Cost Formula

When an AI agent makes a mistake, the total cost of correction includes^[5]:

1. **Detection:** Monitoring systems, user reports, automated anomaly detection
2. **Triage:** Determining severity, scope, and root cause
3. **Intervention:** Human override, rollback, manual correction
4. **Root cause analysis:** Why did this happen?
5. **Prevention:** Updating prompts, retraining, adding guardrails
6. **Verification:** Testing the fix, ensuring it doesn't break other workflows

For complex workflows, correction can require **Tier 3 senior engineers at \$120–\$200/hour**. If a single agent error requires 4 hours of diagnosis + 8 hours of fixing + 4 hours of testing, that is **\$1,920–\$3,200 per incident** at Tier 3 rates.

The break-even calculation: **If an agent processes 10,000 tasks per month at 90% success rate, that is 1,000 failures per month.** If even 10% of those require Tier 3 intervention (100 incidents), the correction cost is \$192K–\$320K monthly — far exceeding the API cost savings.

Reversibility Infrastructure: The Safety Net

To prevent correction costs from spiraling, production agent systems need reversibility^[5]:

- **Transactional memory:** Log the state before the action and after
- **Rollback capabilities:** Undo the agent's action if it was wrong

- **Snapshot backups:** User profiles, database rows, file states
- **Audit trails:** Full record of what changed and why

Building this infrastructure costs **\$30K–\$100K** depending on system complexity^[5]. But the alternative — manual correction at scale — costs orders of magnitude more.

The 89% That Never Launch

According to OneReach AI research cited by multiple sources, **89% of enterprise AI agents never reach production**^[6]. For these projects, the ROI question is not "How much did we save?" but "How much did we spend to learn it won't work?"

Typical sunk costs for a failed agent deployment:

- Initial development: \$100K–\$500K
- Integration work: \$50K–\$200K
- Pilot testing: \$20K–\$80K
- Organizational time (meetings, training, change management): Unquantified but significant

Total write-off: \$170K–\$780K per failed project. The enterprises that succeed are not the ones that avoid failures — they are the ones that fail fast and cheap.

WHAT WOULD INVALIDATE THIS?

If agent reliability improved to 99%+ through architectural breakthroughs (deterministic guardrails, formal verification, etc.), correction costs would drop dramatically. Current reality: most agents operate at 85–95% success rates in production.

SO WHAT?

Budget for failure. Allocate 20% of your agent budget specifically for error correction, rollback infrastructure, and "what if this doesn't work" scenarios. The projects that succeed are the ones that plan for failure from day 1.

8. Real-World Economics: What We Learned Building 18 Reports 95%

(Confidence: High — internal production data)

Ainary has produced 18 research reports using a multi-agent system. Our actual costs provide a reality check against vendor claims.

The Numbers

Production data from Ainary's multi-agent research pipeline (AR-001 through AR-018):

\$2.75

Average API cost per research report (18 reports)

Source: Internal cost tracking, Feb 2026

~40%

Cost reduction through Sonnet vs Opus model switching

Source: Internal A/B comparison

~50%

Token consumption reduction through aggressive caching

Source: Before/after caching implementation

These numbers validate the vendor claim that **per-task API costs are low**. A research report that would take a human analyst 8–16 hours (\$800–\$2,400 at \$100–\$150/hour) costs \$2.75 in API calls.

The Costs Nobody Sees

What the \$2.75 average does not include:

- **System design and architecture:** 40+ hours building the multi-agent pipeline
- **Prompt engineering and optimization:** 20+ hours per report template

- **Quality assurance and fact-checking:** 2–4 hours per report (human oversight)
- **Infrastructure costs:** Cloud hosting, vector database, monitoring
- **Error correction:** Rerunning failed research tasks, fixing prompt injection issues
- **Knowledge management:** Building and maintaining the source quality framework

When the full TCO is calculated, **each report costs \$150–\$300 all-in** (infrastructure + human oversight + API costs). Still cheaper than human-only research (\$800–\$2,400), but **50–100x higher than the API cost alone.**

What Actually Drove Cost Down

The optimizations that mattered:

1. **Model selection:** Using Claude Sonnet 4 for research tasks instead of Opus saved ~40% with minimal quality loss
2. **Aggressive caching:** Caching common research patterns (methodology, source validation) cut token consumption ~50%
3. **Task decomposition:** Breaking reports into 15–20 independent research briefs allowed parallel processing and better error isolation
4. **Structured outputs:** Forcing agents to return JSON reduced post-processing costs
5. **Confidence thresholds:** Auto-rejecting low-confidence outputs prevented expensive human review of garbage

None of these optimizations came from vendor recommendations. All required engineering time to discover and implement.

The Lesson: Economics Are Learnable, But Not Automatic

Our cost per report started at **\$8–\$12** for AR-001 through AR-005. By AR-015, we were at **\$2–\$3**. The cost curve follows a learning function — but you need to invest in the learning.

Enterprises that expect Day 1 costs to match steady-state costs are setting themselves up for disappointment. The economic advantage is real, but it comes from iteration and optimization, not from plugging in an API.

CLAIM

API costs are marginal in mature agent systems, but getting to "mature" requires 10-20x the initial API cost in engineering, testing, and optimization.

WHAT WOULD INVALIDATE THIS?

If plug-and-play agent platforms emerged that delivered optimized performance out-of-the-box with zero tuning, the learning cost would disappear. Current platforms (LangChain, CrewAI, AutoGen) require significant customization for production use.

SO WHAT?

Plan for a 6–12 month learning curve before your agent economics reach steady state. Budget for iteration. Track costs per task from day 1. The teams that optimize aggressively see 3–5x cost reductions over 6 months. The teams that don't are stuck at vendor list prices.

9. When Agent ROI Actually Works 80%

(Confidence: High)

Agent ROI is not a yes/no question — it is a function of volume, complexity, and tolerance for error.

The Break-Even Model

Based on enterprise deployment data, agent ROI becomes positive when^{[1][4][7]}:

1. **Volume threshold:** 50,000+ interactions per year (translates to 4–6 month payback on setup costs)
2. **Task suitability:** Routine, structured, low-stakes tasks where 90% success rate is acceptable
3. **Integration simplicity:** Modern APIs, well-documented systems, minimal legacy infrastructure
4. **Organizational readiness:** Team trained, change management planned, trust infrastructure budgeted

When all four conditions are met, **actual ROI ranges from 3x–10x over 18–24 months**^[12]. When any condition fails, ROI often turns negative.

Exhibit 3: Agent ROI Success Patterns

| USE CASE | VOLUME | SUCCESS RATE NEEDED | TYPICAL ROI |
|---|-----------------------------|--------------------------------------|---|
| Customer service (FAQ, tier 1) | High (100K+/year) | 85–90% | 5–10x (18–24 months) |
| Data entry / form processing | Medium (20K–100K/year) | 95%+ | 3–7x (12–18 months) |
| Content generation (marketing, reports) | Medium (10K–50K tasks/year) | 80–85% (high human review tolerance) | 2–5x (12–24 months) |
| Code generation / review | Variable | 90%+ (security-critical) | 1–3x (24+ months, high correction costs) |
| High-stakes decision-making | Low (<5K/year) | 99%+ (often unachievable) | Negative (trust infrastructure exceeds savings) |

Source: Author synthesis from enterprise case studies [1][7][12] and AR-010 Agent Failure Taxonomy

Where Agents Clearly Win

- **24/7 operations:** Agents don't need shifts, vacations, or overtime pay
- **Extreme scale:** Going from 1 to 1,000 concurrent agents is a config change, not a hiring spree
- **Multi-language support:** Agents handle 50+ languages natively at no extra cost
- **Consistency:** No "bad days," no variation in quality due to fatigue or mood
- **Speed:** Seconds instead of hours for research, analysis, data processing

Where Humans Still Win

- **Judgment under ambiguity:** "Gray zone" decisions where rules don't exist

- **Relationship-building:** High-touch sales, negotiations, executive relationships
- **Creative strategy:** Novel approaches not in training data
- **Accountability:** Legal liability, regulatory compliance, crisis management
- **Empathy:** Customer de-escalation, sensitive conversations, grieving clients

The Hybrid Model: Recommended

The enterprises seeing the best ROI are not replacing humans with agents — they are **augmenting 1 human with 3–5 agents**. The human handles the 20% of tasks that require judgment, relationships, and accountability. The agents handle the 80% that are routine, structured, and automatable.

This model delivers 3–5x productivity gains without the organizational trauma of mass layoffs and without the quality degradation of "agents all the way down."

WHAT WOULD INVALIDATE THIS?

If agents achieved human-level judgment, empathy, and creativity, the hybrid model would become obsolete. Current evidence: agents excel at structured tasks, struggle with ambiguity and novelty.

SO WHAT?

Do not ask "Can agents replace humans?" Ask "Which tasks should agents own, and which should humans own?" The companies that answer this question thoughtfully capture the ROI. The ones that don't end up like Klarna: great cost savings, quality problems, course correction required.

10. Recommendations

Achieving positive agent ROI requires honest cost accounting, aggressive optimization, and organizational readiness. Most teams underinvest in all three.

Scope: These recommendations apply to enterprise teams deploying agents for the first time or struggling to achieve projected ROI. Teams with mature agent deployments may have already learned these lessons.

For Budget Planning

- 1. Budget 1.5x your initial estimate for year 1.** The hidden costs (trust infrastructure, monitoring, security, human oversight) are not optional. Better to plan for them than discover them mid-deployment.
- 2. Allocate 30–50% of dev costs for ongoing maintenance.** Monitoring, retraining, security updates, and infrastructure changes are annual costs, not one-time.
- 3. Budget for failure.** Reserve 20% of your agent budget for error correction, rollback infrastructure, and the possibility that the project fails entirely (89% do).
- 4. Demand TCO transparency from vendors.** Ask: "What costs are not in this quote that we will discover in months 3–12?" The vendors who provide detailed answers are the ones who have deployed at scale.

For Cost Optimization

- 1. Optimize models aggressively.** Start with expensive models (Opus, GPT-4), measure quality, then downgrade to cheaper models (Sonnet, GPT-4o-mini) until quality degrades. We achieved 40% cost savings through Sonnet vs Opus switching with minimal quality loss.
- 2. Implement caching from day 1.** Caching common patterns (prompts, retrieved context, tool outputs) can reduce token consumption 40–50%. This is not a future optimization — it should be in the MVP.

3. **Track cost per task, not cost per API call.** If you measure API costs alone, you will optimize the wrong thing. Measure total cost (API + infra + human oversight) per completed task.
4. **Instrument everything.** Log every agent action, every tool call, every token consumed. The data you collect in months 1–3 will drive the optimizations that matter in months 6–12.

For Organizational Readiness

1. **Start with low-stakes, high-volume tasks.** Customer service tier 1, data entry, content generation. Build trust in agent outputs before deploying to high-stakes workflows.
2. **Plan change management before deployment.** Technical readiness is 30% of success. The other 70% is training, process redesign, and managing employee anxiety. Budget for it.
3. **Build the hybrid model from day 1.** Do not plan to replace humans. Plan to augment them. 1 human + 3–5 agents is the pattern that works.
4. **Fail fast and cheap.** If an agent deployment is not showing positive signals in 3 months, kill it. The sunk cost fallacy is expensive in AI.

SO WHAT?

The ROI is real, but it is not automatic. The enterprises that succeed are the ones that budget honestly, optimize aggressively, and plan for organizational change. The ones that fail are the ones that believe the vendor pitch and allocate for API costs alone.

11. Transparency Note

This section explains the methodology, known limitations, and confidence calibration of this report. We include it because transparency about what we know — and what we don't — is what separates research from marketing.

| | |
|---|---|
| Overall Confidence | 82% |
| Sources | 14 primary (vendor pricing pages, enterprise case studies with disclosed costs, internal production data), 8 secondary (industry surveys, analyst reports) |
| Strongest Evidence | Ainary internal cost data (18 reports, \$2.75 avg API cost, 40% Sonnet savings, 50% caching savings); per-interaction cost advantage (\$0.25–\$0.50 vs \$3–\$8, confirmed by 3 independent sources); maintenance costs 15–30% of dev costs annually (industry-standard estimate) |
| Weakest Point | The "3-7x initial estimate" multiplier is directionally supported but not rigorously quantified across a large sample. Most enterprises do not publicly disclose AI deployment costs, forcing reliance on vendor case studies (selection bias) and practitioner surveys (small samples). |
| What Would Invalidate This Report? | If a large-scale survey ($n > 500$ enterprises) showed that actual costs consistently matched initial estimates, the "hidden cost" thesis would collapse. No such data exists because enterprises don't publish this information. |
| Methodology | Multi-agent research pipeline with cost tracking. Analyzed 14 vendor pricing models, 6 enterprise case studies with disclosed costs, 3 cost estimation frameworks (2025-2026). Compared vendor claims to Ainary's own production data (18 reports). Cross-referenced with AR-002 (Trust Tax), AR-010 (Failure Taxonomy), AR-011 (HITL Illusion), AR-012 (Trust Moat). |
| Limitations | Few enterprises publicly disclose AI agent costs. Most data comes from vendor-published case studies (selection bias toward successful deployments). The "89% never reach |

"production" claim comes from a single vendor (OneReach AI) and is cited widely but not independently verified.

| | |
|--------------------------|---|
| System Disclosure | This report was created with a Multi-Agent Research System with full cost tracking (\$2.89 API cost for this report). |
|--------------------------|---|

12. Claim Register

This register lists the key quantitative and qualitative claims made in this report, with sources and confidence levels. The top 5 claims include explicit invalidation conditions.

Exhibit 4: Claim Register

| # | CLAIM | VALUE | SOURCE | CONFIDENCE |
|----|---|--------------------------|--|--|
| 1 | AI agent cost per interaction | \$0.25–\$0.50 | Teneo.ai [1], Riseup Labs [2], Hypersense [3] | High |
| 2 | Human agent cost per interaction | \$3–\$8 | Teneo.ai [1], Waterfield Tech [7] | High |
| 3 | Maintenance costs (annual % of dev) | 15–30% | Riseup Labs [2] | High (industry standard) |
| 4 | Security/compliance cost adder | 20–40% of platform costs | NoCodeFinder [3] | Medium |
| 5 | Trust infrastructure cost multiplier | 1.5x | Author estimate from AR-002 + case studies [4] | Medium |
| 6 | Enterprise agents never reaching production | 89% | OneReach AI via Hypersense [6] | Low (single source, methodology unclear) |
| 7 | Correction cost break-even threshold | 10% error rate | Leena.AI [5] | Medium |
| 8 | Klarna AI savings | \$60M, 853 FTEs replaced | CX Dive [12], The Guardian, Yahoo Finance | High (multiple independent sources) |
| 9 | Klarna "overpivot" on AI | CEO admission | Forbes Jan 2026 [13] | High |
| 10 | Ainary avg report API cost | \$2.75 | Internal cost tracking (18 reports) | High (direct measurement) |

| | | | | |
|----|-----------------------------|------|-------------------------|---------------------------|
| 11 | Sonnet vs Opus cost savings | ~40% | Internal A/B comparison | High (direct measurement) |
| 12 | Caching token reduction | ~50% | Internal before/after | High (direct measurement) |

Top 5 Claims — Invalidation Conditions:

- **Claim #1 & #2 (per-interaction costs):** Invalidated if large-scale independent audit ($n > 1,000$ enterprises) shows actual per-interaction costs differ by $>50\%$ from these ranges.
- **Claim #3 (maintenance 15–30%):** Invalidated if industry-standard accounting frameworks shift or if zero-maintenance agent platforms emerge.
- **Claim #5 (1.5x trust multiplier):** Invalidated if rigorous TCO study across >100 deployments shows trust infrastructure costs are $<1.2x$ or $>2.0x$.
- **Claim #6 (89% never deploy):** Invalidated if independent research with disclosed methodology shows deployment success rates $>50\%$. This claim has low confidence and should be treated as directional.

13. References

- [1] Teneo.ai. (2025). "AI vs Live Agent Cost: The Complete 2025 Analysis and Comparison."
- [2] Riseup Labs. (2025). "AI Agent Development Cost: Full Breakdown for 2026."
- [3] Hypersense Software. (2026). "The Hidden Costs of AI Agent Development: A Complete TCO Guide for 2026."
- [4] NoCodeFinder. (2025). "AI Agent Pricing 2026: Complete Cost Guide & Calculator."
- [5] Leena.AI. (2026). "Automation and AI: Measuring the Correction Cost of Errors."
- [6] OneReach AI via Hypersense Software. "89% of enterprise AI agents never reach production."
- [7] Agent Framework Hub / MintSquare. (2026). "AI Agent Production Costs 2026: Real Data."
- [8] Product Crafters. (2026). "AI Agent Development Cost: \$5K to \$500K+ (2026 Pricing)."
- [9] Microsoft / NVD. (2025). "CVE-2025-32711: EchoLeak — Information Disclosure in Microsoft Copilot."
- [10] Author analysis based on AR-016 (Agent Observability Gap) and vendor pricing for LangSmith, Arize, Langfuse.
- [11] Symphonize. "Costs of Building AI Agents: What Decision Makers Need to Know." Cites McKinsey/BCG research on change management costs exceeding technology costs.
- [12] CX Dive. (2025). "Klarna says AI agent does work of 853 employees."
- [13] Forbes. (2026). "AI Productivity's \$4 Trillion Question: Hype, Hope, and Hard Data." CEO Siemiatkowski admits "overpivoted."

Cite as: Ainary Research (2026). *AI Agent Economics — The Real ROI Nobody Talks About*. AR-021.

About the Author

Florian Ziesche is the founder of Ainary Ventures, where AI does 80% of the research and humans do the 20% that matters. Before Ainary, he was CEO of 36ZERO Vision and advised startups and SMEs on AI strategy and due diligence. His conviction: HUMAN × AI = LEVERAGE. This report is the proof.

ainaryventures.com



AI Strategy · Published Research · Daily Intelligence

Contact · Feedback

ainaryventures.com

florian@ainaryventures.com

© 2026 Ainary Ventures